HSI Special Period Model Performance Analysis

> ***Experiment***: *SpecialPeriod_Seed2620*
> ***Market***: *Hang Seng Index (^HSI)*
> ***Models***: *ARIMA, Bi-RNN, Bi-GRU, Bi-LSTM, PatchTST, Mamba, SNN, Ensemble*
> ***Window Sizes***: *30, 60, 90, 120, 252*
> ***Prediction Horizons***: *1-day, 5-day, 10-day*
> ***Evaluation Metrics***: *R^2, RMSE, MAE*

---

1. Overall Performance Summary

1.1 Period Difficulty Ranking (Average R^2 across all models)

Rank	Period	Avg R^2	Avg RMSE	Avg MAE	Samples	Difficulty
1	2008 Financial Crisis	**0.8335**	0.0812	0.0661	367	Easiest
2	2022 Rate Hike	0.7961	0.0392	0.0309	246	
3	2015 China Crash	0.7961	0.0377	0.0295	186	
4	2010 Euro Debt Crisis	0.7709	0.0319	0.0253	555	
5	2018 Trade War	0.7582	0.0394	0.0323	246	
6	2023 AI Wave	0.7079	0.0280	0.0224	243	
7	2020 COVID Crash	**0.3237**	0.0433	0.0322	102	Hardest

Interpretation:

- **2008 Financial Crisis (R^2=0.83)**: The HSI experienced a dramatic decline from its October 2007 peak (~31,600) to its March 2009 trough (~11,300), a 64% drop that was even more severe than the SP500's 57% decline. However, this sharp, persistent downtrend with high autocorrelation made it highly predictable for time-series models. The large sample size (367 days) further aided model learning.
- **2022 Rate Hike & 2015 China Crash (both R^2=0.80)**: These two periods tied for second place. The 2015 China Crash is particularly notable — while SP500 found this period difficult (R^2=0.21), HSI performed excellently (R^2=0.80). This is because the HSI was directly affected by the Chinese market crash rather than experiencing indirect spillover, making the price dynamics more structured and predictable. The 2022 Rate Hike also impacted HSI more directly through Hong Kong's USD peg, creating clear monetary policy transmission channels.

- **2010 Euro Debt Crisis (R^2=0.77)**: With the largest sample size (555 days) and moderate volatility, this period provided ample training data. HSI's exposure to European events was primarily through trade channels, creating a more gradual and predictable impact pattern.
- **2018 Trade War (R^2=0.76)**: Unlike SP500 (R^2=0.20), HSI was a direct participant in the US-China trade war, making the market's reactions more consistent and predictable from historical patterns rather than relying on unpredictable tweet-driven signals.
- **2023 AI Wave (R^2=0.71)**: HSI benefited less from the AI boom compared to the tech-heavy SP500, resulting in less volatile but still directional movements that models could capture reasonably well.
- **2020 COVID Crash (R^2=0.32)**: As with SP500, the pandemic's unprecedented nature made this the most challenging period. However, HSI's R^2 (0.32) was still significantly better than SP500's (0.07), partly because Hong Kong's market response was somewhat tempered by proximity to China's earlier recovery trajectory.

1.2 Model Overall Ranking (Average R^2 across all conditions)

Rank	Model	Avg R^2	Avg RMSE	Avg MAE	Std(R^2)	Stability
1	**ARIMA**	**0.8188**	0.0329	0.0256	0.1421	Most Stable
2	Bi-GRU	0.8090	0.0343	0.0272	0.1556	
3	Bi-RNN	0.8071	0.0342	0.0269	0.1551	Stable
4	Mamba	0.7996	0.0361	0.0283	0.1572	
5	Bi-LSTM	0.7984	0.0360	0.0283	0.1540	
6	Ensemble	0.7890	0.0379	0.0295	0.1559	
7	PatchTST	0.5182	0.0627	0.0511	0.2229	
8	SNN	0.3585	0.0694	0.0561	0.3325	Least Stable

Interpretation:

- **Remarkably tight top-6 cluster**: Unlike SP500 where models ranged from -0.80 to 0.75 in R^2, the HSI's top 6 models are clustered within a narrow band of 0.79–0.82. This suggests that HSI's price dynamics are more uniformly capturable across different architectures, likely due to stronger autocorrelation and more structured market microstructure.
- **ARIMA leads but by a thin margin**: ARIMA's advantage over Bi-GRU is only 0.01 R^2, compared to the 0.11 gap on SP500. This indicates that deep learning models are much more competitive on HSI, possibly because HSI exhibits nonlinear patterns (e.g., from China-HK market interactions) that deep learning can exploit.
- **Bi-GRU edges out Bi-RNN**: On HSI, Bi-GRU (0.809) slightly outperforms Bi-RNN (0.807), reversing the SP500 pattern where Bi-RNN was superior. GRU's gating mechanism may be better at capturing the cross-market interaction patterns characteristic of HSI.
- **PatchTST achieves positive R^2 (0.52)**: A dramatic improvement over SP500 (-0.80). HSI's stronger trend persistence and higher autocorrelation may provide the long-range patterns that PatchTST's patch-based attention mechanism is designed to capture.

- **SNN remains the weakest but still positive (0.36)**: Even SNN achieves positive R^2 on HSI (vs -0.13 on SP500), further confirming that HSI is inherently more predictable.
- **All models show similar stability**: The std(R^2) values for the top 6 models are remarkably similar (0.14–0.16), indicating that HSI's predictability is consistently higher across market regimes regardless of model architecture.

---

2. Analysis by Prediction Horizon

2.1 Horizon Effect on R^2

Horizon	Avg R^2	Avg RMSE	Avg MAE
1-day	**0.8725**	0.0281	0.0221
5-day	0.7283	0.0444	0.0354
10-day	0.5361	0.0564	0.0449

2.2 R^2 by Period x Horizon

Period	H=1	H=5	H=10	Degradation (H=1→H=10)
2015 China Crash	0.9189	0.8204	0.6488	-29.4%
2022 Rate Hike	0.9181	0.8126	0.6576	-28.4%
2008 Financial Crisis	0.8998	0.8265	0.7742	-14.0%
2010 Euro Debt Crisis	0.8842	0.7865	0.6420	-27.4%
2018 Trade War	0.8724	0.7485	0.6538	-25.1%
2023 AI Wave	0.8553	0.6817	0.5866	-31.4%
2020 COVID Crash	0.7590	0.4221	**-0.2102**	-127.7% (collapse)

Interpretation:

- **HSI is far more horizon-resilient than SP500**: R^2 at H=10 remains 0.54 (vs SP500's 0.10), and the degradation from H=1 to H=10 is only 39% (vs SP500's 83%). This suggests that HSI's price trends persist more strongly over multi-day horizons.
- **2008 Crisis shows the least degradation (-14%)**: HSI's steep, persistent downtrend during the financial crisis maintained strong autocorrelation even at 10-day horizons, allowing models to sustain R^2=0.77 at H=10.
- **Six of seven periods maintain positive R^2 at H=10**: Unlike SP500 where 3 periods had negative R^2 at H=10, only the 2020 COVID period collapses to negative territory on HSI. This demonstrates HSI's superior long-horizon predictability.

- **2020 COVID is the sole exception**: The COVID crash remains the only period where models fail at H=10 (R^2=-0.21), though even here HSI performs substantially better than SP500 (which reached -0.47).

2.3 Best Deep Learning Model per Period x Horizon

Period	H=1	H=5	H=10
2008 Financial Crisis	Bi-RNN (0.9826)	Bi-RNN (0.9350)	Bi-LSTM (0.8929)
2010 Euro Debt Crisis	Bi-RNN (0.9727)	Bi-RNN (0.8734)	Bi-RNN (0.7658)
2015 China Crash	Bi-GRU (0.9781)	Bi-LSTM (0.8869)	Bi-GRU (0.7298)
2018 Trade War	Bi-GRU (0.9707)	Bi-GRU (0.8588)	Bi-GRU (0.7820)
2020 COVID Crash	Bi-LSTM (0.9108)	Mamba (0.5724)	Bi-RNN (-0.0785)
2022 Rate Hike	Bi-RNN (0.9697)	Bi-RNN (0.8584)	Mamba (0.7068)
2023 AI Wave	Mamba (0.9670)	Bi-GRU (0.8534)	Mamba (0.7380)

Interpretation:

- **Bi-RNN and Bi-GRU dominate across horizons**: Unlike SP500 where Mamba dominated H=1, HSI shows a preference for simpler recurrent models. Bi-RNN leads at H=1 in 3/7 periods and Bi-GRU leads in 2/7, while Mamba only leads in 1/7.
- **Model leadership is more stable across horizons**: On SP500, different models dominated at different horizons (Mamba at H=1, Bi-RNN at H=10). On HSI, the same model often leads across multiple horizons (e.g., Bi-RNN leads at both H=1 and H=5 for 2008, 2010, and 2022).
- **Mamba excels in recent periods (2022–2023)**: Mamba's state-space mechanism appears well-suited for capturing the more complex cross-market dynamics of recent years.
- **Deep learning achieves >0.97 at H=1**: The best DL models at H=1 are extremely close to ARIMA (gap < 0.01), confirming that on HSI, DL models match statistical baselines for short-term prediction.

---

3. Analysis by Window Size

3.1 Window Size Effect on R^2

Window	Avg R^2	Avg RMSE
30	**0.7660**	0.0387
60	0.7333	0.0411
90	0.6979	0.0446
120	0.7054	0.0442
252	0.6590	0.0461

3.2 R^2 by Window x Horizon (averaged across all models and periods)

Window \ Horizon	H=1	H=5	H=10
30	**0.9157**	**0.7896**	**0.5927**
60	0.8950	0.7524	0.5525
90	0.8687	0.7183	0.5067
120	0.8644	0.7292	0.5224
252	0.8189	0.6522	0.5061

Interpretation:

- **Shorter windows are clearly better on HSI**: W=30 (R^2=0.77) significantly outperforms W=252 (R^2=0.66), a 0.11 gap that is larger than the window effect on SP500 (0.12). This is the opposite of ARIMA's preference on SP500, where larger windows helped.
- **Performance decreases monotonically with window size**: Unlike SP500's non-monotonic pattern (where W=120 was optimal), HSI shows a clear trend: more recent data is always better. This suggests HSI's market dynamics change more rapidly, making older data actively harmful.
- **W=30 dominates across all horizons**: At every horizon (H=1, H=5, H=10), W=30 is the best window size. The advantage is most pronounced at H=5 (0.79 vs 0.65 for W=252).
- **Possible explanation**: HSI is more sensitive to rapidly changing China-HK policy environments, cross-border capital flows, and regional geopolitical events. A 30-day window captures the current regime without contamination from previous, potentially very different, market conditions.

3.3 Window Size Effect per Model

Model \ Window	W=30	W=60	W=90	W=120	W=252
ARIMA	0.8027	0.8231	0.8171	0.8237	0.8298
Bi-RNN	0.8132	0.8070	0.8009	0.8114	0.8032
Bi-GRU	0.8127	0.8095	0.8109	0.8024	0.8094
Bi-LSTM	0.8044	0.7938	0.7985	0.8032	0.7889
Mamba	0.8077	0.7970	0.7923	0.7990	0.8017
PatchTST	0.6440	0.6545	0.4828	0.5130	0.2301
SNN	0.6745	0.3652	0.2724	0.2706	0.1531
Ensemble	0.8018	0.7879	0.7764	0.7903	0.7732

Interpretation:

- **ARIMA is the only model that prefers larger windows**: ARIMA's performance improves monotonically from W=30 (0.80) to W=252 (0.83), consistent with its need for more data to estimate statistical parameters.
- **Deep learning models are remarkably window-insensitive**: Bi-RNN, Bi-GRU, Bi-LSTM, and Mamba all maintain R^2 within a tight 0.79–0.81 band regardless of window size. This is a stark contrast to SP500 where Bi-RNN dropped from 0.68 (W=30) to 0.47 (W=252).
- **PatchTST and SNN are highly window-sensitive**: PatchTST drops from 0.64 (W=30) to 0.23 (W=252), and SNN from 0.67 to 0.15. These models cannot handle the noise introduced by longer lookback periods.

---

4. Detailed Period-by-Period Analysis

4.1 2008 Financial Crisis (2007-10 to 2009-03)

Model	H=1	H=5	H=10	Avg
ARIMA	0.9884	0.9562	0.9194	**0.9547**
Bi-RNN	0.9826	0.9350	0.8589	0.9255
Bi-GRU	0.9809	0.9151	0.8557	0.9172
Bi-LSTM	0.9552	0.8781	0.8929	0.9087
Mamba	0.9643	0.8980	0.8088	0.8903
Ensemble	0.9572	0.8838	0.8371	0.8927
SNN	0.7370	0.6693	0.5543	0.6535
PatchTST	0.6326	0.4764	0.4667	0.5252

Analysis: HSI's 2008 crisis performance is exceptional across all models — even PatchTST achieves a positive R^2 of 0.53, and SNN reaches 0.65. The HSI experienced a more severe drawdown than SP500 (64% vs 57%), but the decline was extremely persistent and directional, creating an ideal signal for time-series models. All four mainstream DL models (Bi-RNN, Bi-GRU, Bi-LSTM, Mamba) achieve $R^2 > 0.89$, demonstrating that the crisis's trend-driven dynamics were universally capturable. The R^2 at H=10 remains remarkably high (0.77 average), reflecting the sustained momentum of the crash.

4.2 2010 Euro Debt Crisis (2010-04 to 2012-06)

Model	H=1	H=5	H=10	Avg
ARIMA	0.9756	0.8749	0.7539	**0.8681**
Bi-RNN	0.9727	0.8734	0.7658	0.8706
Bi-GRU	0.9722	0.8725	0.7548	0.8665
Mamba	0.9718	0.8587	0.7451	0.8585
Bi-LSTM	0.9702	0.8599	0.7324	0.8542
Ensemble	0.9407	0.8516	0.7414	0.8446
PatchTST	0.6635	0.6094	0.3655	0.5462
SNN	0.6065	0.4916	0.2770	0.4584

Analysis: The Euro Debt Crisis shows remarkably uniform performance among the top 5 models (R^2 range: 0.85–0.87). Bi-RNN slightly outperforms ARIMA (0.8706 vs 0.8681), making this one of the rare periods where a DL model edges out the statistical baseline. The large sample size (555 days) and moderate volatility (lowest RMSE=0.032) created ideal conditions. At H=1, all top 5 models achieve >0.97, indicating near-perfect next-day prediction. The consistent performance across models suggests that the Euro Debt Crisis's impact on HSI was transmitted through well-structured channels (trade flows, capital movement) that multiple architectures could learn equally well.

4.3 2015 China Crash (2015-06 to 2016-02)

Model	H=1	H=5	H=10	Avg
ARIMA	0.9771	0.8865	0.7649	**0.8762**
Bi-GRU	0.9781	0.8824	0.7298	0.8634
Mamba	0.9772	0.8776	0.7162	0.8570
Bi-LSTM	0.9778	0.8869	0.7023	0.8557
Bi-RNN	0.9739	0.8784	0.7106	0.8543
Ensemble	0.9572	0.8612	0.6982	0.8389
SNN	0.6552	0.6369	0.5558	0.6160
PatchTST	0.8551	0.6535	0.3129	0.6072

Analysis: This is one of the most striking HSI-SP500 divergences. While SP500 only achieved R^2=0.21 for this period, HSI reaches R^2=0.80. The reason is clear: the 2015 China stock market crash directly affected the Hong Kong market through the Stock Connect program and shared investor base, making HSI's reaction structured and predictable from its own historical data. In contrast, the SP500's indirect spillover through contagion channels was noisy and difficult to model. PatchTST achieves an impressive 0.86 at H=1, suggesting that the crash's patch-level patterns were particularly well-suited to its architecture.

4.4 2018 Trade War (2018-01 to 2018-12)

Model	H=1	H=5	H=10	Avg
Bi-GRU	0.9707	0.8588	0.7820	**0.8705**
Ensemble	0.9472	0.8548	0.7847	0.8622
Bi-RNN	0.9540	0.8454	0.7655	0.8550
Mamba	0.9677	0.8381	0.7459	0.8506
ARIMA	0.9719	0.8473	0.7204	0.8465
Bi-LSTM	0.9658	0.8405	0.7248	0.8437
PatchTST	0.8161	0.7651	0.6210	0.7341
SNN	0.3856	0.1380	0.0857	0.2031

Analysis: A remarkable result — **Bi-GRU (0.87) outperforms ARIMA (0.85)**, making the 2018 Trade War one of the few periods where a DL model is the overall best performer. This is a dramatic contrast with SP500, where R^2=0.20. HSI was a direct battleground of the US-China trade war, meaning its price reactions to trade-related events were more immediate, consistent, and learnable. The Ensemble also performs exceptionally well (0.86), second only to Bi-GRU, suggesting that aggregating multiple perspectives helps in this complex environment. PatchTST achieves its second-best performance (0.73), indicating that trade-war patterns had structured temporal features. SNN, however, struggles severely (0.20), suggesting it cannot model the conditional dynamics of trade negotiations.

4.5 2020 COVID Crash (2020-02 to 2020-06)

Model	H=1	H=5	H=10	Avg
ARIMA	0.9043	0.5530	0.0723	**0.5099**
Bi-RNN	0.9046	0.5587	-0.0785	0.4616
Bi-GRU	0.9089	0.5615	-0.0883	0.4607
Bi-LSTM	0.9108	0.5655	-0.1148	0.4538
Mamba	0.9080	0.5724	-0.1464	0.4446
Ensemble	0.8570	0.5344	-0.0713	0.4400
PatchTST	0.5857	0.1632	-0.5606	0.0628
SNN	0.0929	-0.1318	-0.6938	-0.2442

Analysis: The COVID crash remains the most challenging period, though HSI handles it substantially better than SP500 (avg R^2=0.32 vs 0.07). At H=1, all mainstream models achieve R^2>0.90 — even during the pandemic, next-day prediction remains reliable. The critical breakdown occurs at H=10, where all DL models turn negative. Interestingly, ARIMA is the only model that maintains positive R^2 at H=10 (0.07), suggesting that even weak autocorrelation structure is more reliable than neural network extrapolation during unprecedented events. The fact that Bi-LSTM (0.91) marginally outperforms ARIMA (0.90) at H=1 during COVID is notable — LSTM's memory gates may help during rapid regime changes by selectively retaining or discarding pre-crisis patterns.

4.6 2022 Rate Hike (2022-01 to 2022-12)

Model	H=1	H=5	H=10	Avg
Mamba	0.9688	0.8556	0.7068	**0.8437**
ARIMA	0.9684	0.8528	0.7060	0.8424
Bi-RNN	0.9697	0.8584	0.6991	0.8424
Bi-GRU	0.9693	0.8546	0.7025	0.8421
Bi-LSTM	0.9676	0.8514	0.6727	0.8306
Ensemble	0.9511	0.8446	0.6916	0.8291
PatchTST	0.8047	0.7284	0.5541	0.6957
SNN	0.7449	0.6548	0.5283	0.6427

Analysis: The 2022 rate hike period shows the tightest model clustering of all periods — the top 5 models all fall within R^2 0.83–0.84. Mamba edges out the competition (0.8437), consistent with its strong performance on SP500's 2022 as well. Hong Kong's currency peg to the USD means that Fed rate hikes directly impact HK monetary conditions, creating a highly structured, policy-driven environment. Even PatchTST (0.70) and SNN (0.64) perform well here, confirming that the 2022 rate hike created

clear, learnable patterns. At H=10, all DL models maintain R^2 above 0.67, demonstrating that policy-driven trends are persistent enough for medium-term prediction.

4.7 2023 AI Wave (2023-01 to 2023-12)

Model	H=1	H=5	H=10	Avg
Mamba	0.9670	0.8512	0.7380	**0.8521**
Bi-GRU	0.9665	0.8534	0.7070	0.8423
Bi-LSTM	0.9664	0.8515	0.7089	0.8423
Bi-RNN	0.9658	0.8473	0.7089	0.8406
ARIMA	0.9648	0.8415	0.6945	0.8336
Ensemble	0.9348	0.8105	0.7013	0.8155
PatchTST	0.7792	0.2103	0.3797	0.4564
SNN	0.2982	0.1877	0.0542	0.1801

Analysis: Another period where **deep learning outperforms ARIMA** — Mamba (0.85), Bi-GRU (0.84), Bi-LSTM (0.84), and Bi-RNN (0.84) all surpass ARIMA (0.83). This may be because the 2023 AI wave created complex, nonlinear market dynamics that DL models can capture but ARIMA's linear autoregressive structure cannot. Mamba's state-space architecture appears particularly well-suited for the 2023 market, achieving the highest R^2 at both H=1 (0.967) and H=10 (0.738). HSI's response to the AI wave was more muted than SP500's (which had direct exposure to NVIDIA, Microsoft, etc.), creating less noisy but still directional patterns.

---

5. HSI vs SP500 Comparative Analysis

5.1 Overall Comparison

Dimension	SP500	HSI	Difference
Overall Avg R^2	0.34	**0.71**	+0.37
Best Model R^2	0.75 (ARIMA)	**0.82 (ARIMA)**	+0.07
Worst Model R^2	-0.80 (PatchTST)	**0.36 (SNN)**	+1.16
H=1 R^2	0.57	**0.87**	+0.30
H=10 R^2	0.10	**0.54**	+0.44
Model Std Range	0.16–0.74	**0.14–0.33**	Much tighter

5.2 R^2 Comparison by Model

Model	SP500 R^2	HSI R^2	Improvement
ARIMA	0.7501	0.8188	+0.07
Bi-GRU	0.5827	0.8090	**+0.23**
Bi-RNN	0.6303	0.8071	+0.18
Mamba	0.6367	0.7996	+0.16
Bi-LSTM	0.4798	0.7984	**+0.32**
Ensemble	0.5687	0.7890	+0.22
PatchTST	-0.7999	0.5182	**+1.32**
SNN	-0.1282	0.3585	+0.49

5.3 R^2 Comparison by Period

Period	SP500 R^2	HSI R^2	Improvement
2015 China Crash	0.2145	0.7961	**+0.58**
2018 Trade War	0.1983	0.7582	**+0.56**
2022 Rate Hike	0.2723	0.7961	**+0.52**
2023 AI Wave	0.3604	0.7079	+0.35
2020 COVID Crash	0.0683	0.3237	+0.26
2010 Euro Debt Crisis	0.5395	0.7709	+0.23
2008 Financial Crisis	0.7268	0.8335	+0.11

Interpretation:
HSI consistently outperforms SP500 across every model, every period, and every horizon. The most striking improvements are:

1. **PatchTST improvement (+1.32 R^2)**: From completely broken on SP500 (-0.80) to reasonably functional on HSI (0.52). This strongly suggests PatchTST's failure is market-specific, not architectural — HSI's stronger autocorrelation provides the long-range patterns that PatchTST needs.
2. **2015 China Crash improvement (+0.58 R^2)**: The largest period-level improvement, perfectly explained by market geography — HSI is directly exposed to China's market dynamics while SP500 only experiences indirect spillover.
3. **2018 Trade War improvement (+0.56 R^2)**: Similarly explained by direct participation — HSI was a primary battleground of the trade war.
4. **Deep learning models benefit more than ARIMA**: ARIMA improves by +0.07, while DL models improve by +0.16 to +0.32. This suggests that HSI has stronger nonlinear patterns that DL can exploit but ARIMA cannot.

5.4 Why HSI is More Predictable

Several structural factors explain HSI's superior predictability:

1. **Stronger autocorrelation**: HSI exhibits higher serial correlation due to the influence of the Chinese market, which has documented inefficiencies from retail investor dominance, regulatory interventions, and capital controls.
2. **Direct exposure to regional events**: When analyzing events like the China Crash or Trade War, HSI was a direct participant rather than an indirect recipient of spillover. Direct exposure creates more structured, learnable price patterns.
3. **Market microstructure**: HSI's trading hours overlap with multiple Asian markets (Shanghai, Shenzhen, Tokyo), and the opening price often reflects overnight information from US markets, creating predictable "information catch-up" patterns.
4. **Currency peg dynamics**: Hong Kong's USD peg means monetary policy is imported directly from the Fed, creating mechanical relationships between US rates and HK asset prices.
5. **Lower market efficiency**: Compared to the highly efficient SP500, HSI has a higher proportion of retail investors and is subject to periodic regulatory interventions from both HK and mainland Chinese authorities, creating exploitable patterns.

---

6. Cross-Cutting Insights

6.1 Trend-Driven vs. Event-Driven Periods

Category	Periods	HSI Avg R^2	SP500 Avg R^2
Trend-driven	2008 Crisis, 2010 Euro Debt, 2023 AI Wave	**0.77**	0.54
Event-driven	2015 China Crash, 2018 Trade War, 2020 COVID	**0.63**	0.16
Mixed	2022 Rate Hike	**0.80**	0.27

Key difference: HSI's "event-driven" periods are still highly predictable (R^2=0.63 vs SP500's 0.16) because the events directly affected HSI rather than being exogenous shocks.

6.2 The Horizon-Complexity Trade-off

Model Complexity	H=1	H=5	H=10
Low (ARIMA, Bi-RNN)	0.97	0.86	0.71
Medium (Bi-GRU, Mamba)	0.97	0.85	0.71
High (PatchTST, SNN)	0.60	0.44	0.24

On HSI, the low and medium complexity models perform identically, suggesting that the bias-variance trade-off is less critical when the underlying signal is strong.

6.3 Model Architecture Suitability Summary

Model	Strengths on HSI	HSI vs SP500	Best Use Case
ARIMA	Highest overall R^2, most stable	Moderate gain (+0.07)	Universal baseline
Bi-GRU	Best in Trade War, strong overall	Large gain (+0.23)	Event-driven periods
Bi-RNN	Best in 2008/2010, very stable	Large gain (+0.18)	Crisis periods, long horizons
Mamba	Best in 2022/2023, excels at H=10	Large gain (+0.16)	Recent/policy-driven periods
Bi-LSTM	Competitive across all periods	Largest DL gain (+0.32)	General purpose on HSI
Ensemble	Consistent, benefits from strong components	Large gain (+0.22)	Risk-averse deployment
PatchTST	Positive R^2 on HSI (0.52)	Dramatic gain (+1.32)	Only on high-autocorrelation markets
SNN	Positive R^2 on HSI (0.36)	Large gain (+0.49)	Not recommended despite improvement

6.4 Key Takeaways

1. **HSI is substantially more predictable than SP500**: Average R^2 of 0.71 vs 0.34. This improvement holds across all models, periods, and horizons, suggesting it is a fundamental market characteristic rather than a model artifact.
2. **Deep learning models nearly match ARIMA on HSI**: The gap between ARIMA (0.82) and the best DL model (0.81) is negligible. In two periods (2018 Trade War, 2023 AI Wave), DL models actually outperform ARIMA, suggesting that HSI's nonlinear dynamics reward neural network architectures.
3. **HSI maintains predictability at longer horizons**: R^2 at H=10 is 0.54 (vs SP500's 0.10). This has significant practical implications — medium-term trading strategies may be viable on HSI where they would fail on SP500.
4. **Market geography matters more than model architecture**: The largest performance gains come from the HSI market itself (e.g., 2015 China Crash: +0.58 R^2), dwarfing any model selection effect. This underscores the importance of market selection in quantitative finance research.
5. **Even weak models work on HSI**: PatchTST achieves positive R^2 (0.52) and SNN achieves 0.36, compared to catastrophic failure on SP500. This confirms that HSI's signal-to-noise ratio is fundamentally higher.

Report generated from SpecialPeriod_Seed2620 experiment data.